

prithviraj

Intelligent Document Classifier and Router - Product Requirements Document

Try HackMD

Intelligent Document Classifier and Router - Product Requirements Document

1. Executive Summary

1.1 Product Vision

Eliminate the productivity drain of manual document management by intelligently classifying, routing, and

1.2 Business Objectives

Productivity: Save 1 full workday per week per knowledge worker

Speed: 70% reduction in document routing time (3 days → 8 hours)

Accuracy: 95% correct document classification rate

Searchability: 80% reduction in time spent searching for documents

1.3 Target Market

Knowledge-intensive businesses (legal, consulting, finance)

Companies with document-heavy workflows

Organizations with compliance requirements

Remote/hybrid teams needing document collaboration

2. Product Overview

2.1 Core Value Proposition

AI-powered document intelligence that automatically understands, categorizes, and routes business docum

2.2 Key Features

Intelligent Classification: AI-powered document type and content analysis

Smart Routing: Context-aware assignment to teams and individuals

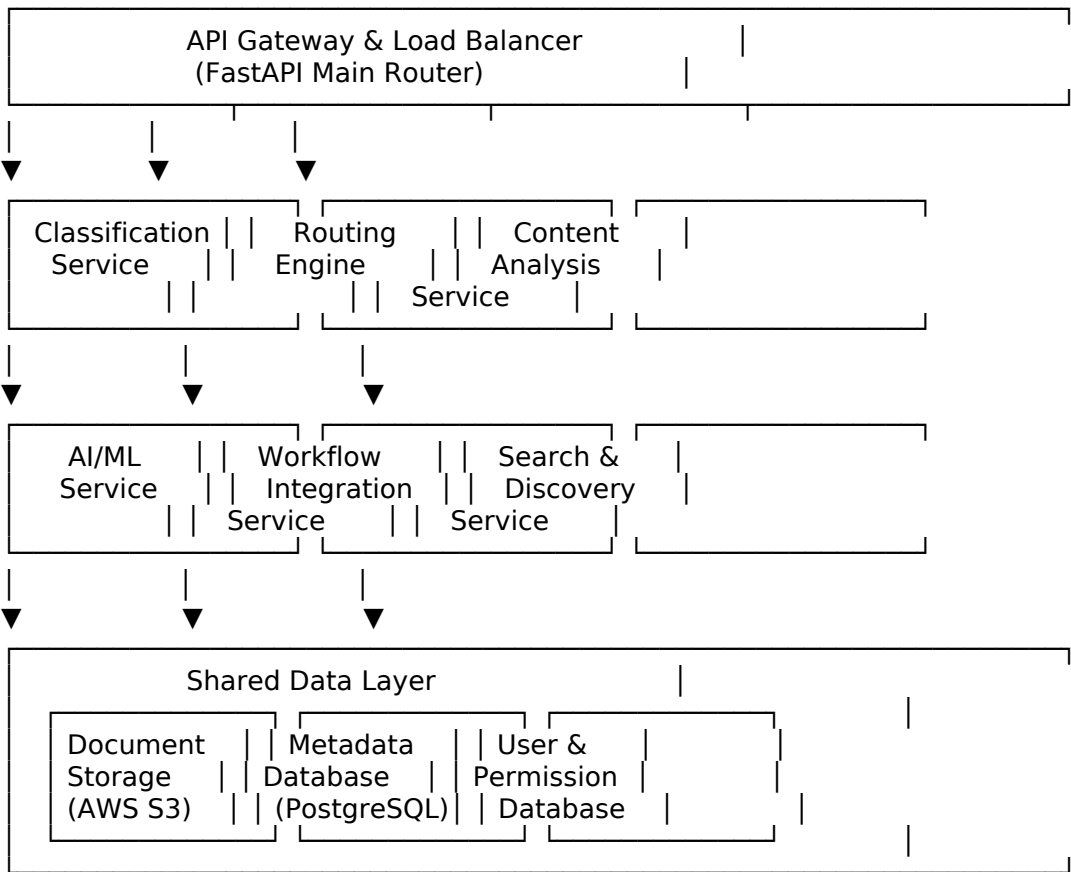
Content Extraction: Key information extraction and tagging

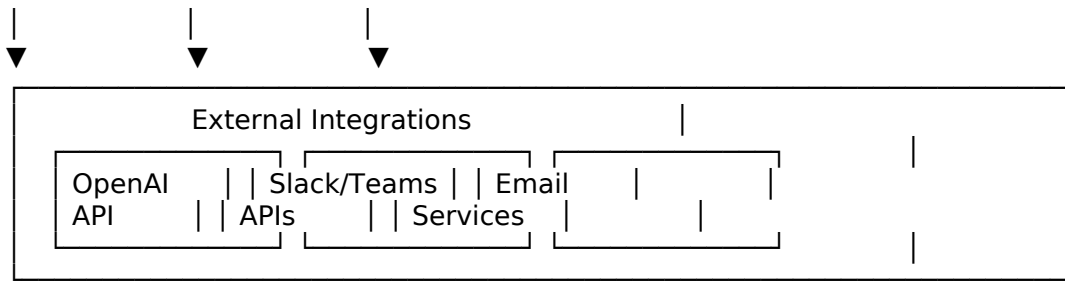
Workflow Integration: Seamless integration with existing business tools

Knowledge Discovery: AI-powered insights and document relationships

3. Microservice Architecture

3.1 High-Level Architecture Diagram





3.2 Detailed Service Architecture

Classification Service

- Document Ingestion
 - Multi-format support (PDF, DOC, TXT, images)
 - Batch processing
 - Real-time processing
- Content Analysis
 - Text extraction
 - OCR processing
 - Metadata extraction
- AI Classification Engine
 - Document type detection
 - Content categorization
 - Confidence scoring
- Learning Module
- User feedback integration
- Model fine-tuning
- Performance optimization

Routing Engine

- Rule Engine
 - Business rule evaluation
 - Priority assignment
 - Escalation logic
- Assignment Logic
 - Team/individual routing
 - Workload balancing
 - Expertise matching
- Context Analysis
 - Project association
 - Customer relationship
 - Historical patterns
- Decision Tracking
- Routing decisions log
- Performance metrics
- Optimization insights

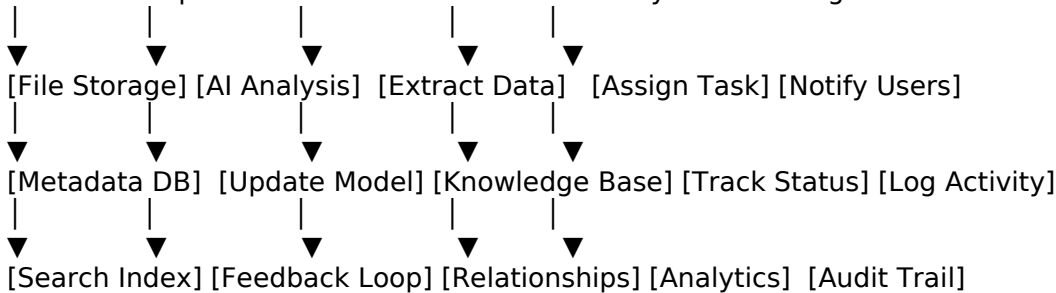
Content Analysis Service

- Information Extraction
 - Entity recognition (names, dates, amounts)
 - Key phrase extraction
 - Summary generation
- Relationship Mapping
 - Document connections
 - Reference detection
 - Dependency analysis
- Compliance Scanning
 - Sensitive data detection
 - Regulatory compliance
 - Risk assessment
- Insight Generation
- Actionable insights

└ Trend analysis
└ Recommendation engine

3.3 Data Flow Diagram

Document Input → Classification → Content Analysis → Routing → Notification



4. Feature Specifications

4.1 Document Classification Features

4.1.1 Multi-Format Support

Document Types: PDF, Word, Excel, PowerPoint, images, text files

Email Integration: Direct email attachment processing

Drag & Drop: Browser-based file upload interface

Batch Processing: Multiple file handling

API Upload: Programmatic document submission

4.1.2 AI-Powered Classification

Document Types: Contracts, invoices, reports, correspondence, legal docs

Content Categories: Financial, legal, technical, marketing, HR

Industry-Specific: Healthcare, finance, legal, manufacturing classifications

Custom Categories: User-defined classification schemes

Confidence Levels: AI certainty scores for classifications

4.1.3 Content Understanding

Entity Extraction: People, organizations, dates, amounts, locations

Key Information: Critical data points and summaries

Language Detection: Multi-language document support

Sentiment Analysis: Document tone and sentiment scoring

Topic Modeling: Automatic subject identification

4.2 Intelligent Routing Features

4.2.1 Smart Assignment

Role-Based Routing: Route based on user roles and responsibilities

Expertise Matching: Match documents to subject matter experts

Workload Balancing: Distribute work evenly across team members

Priority-Based: High-priority documents to appropriate resources

Time Zone Awareness: Route to available team members

4.2.2 Business Rule Engine

Conditional Routing: If-then-else routing logic

Multi-Criteria: Route based on multiple document attributes

Escalation Rules: Automatic escalation for unhandled documents

Exception Handling: Special handling for edge cases

Approval Workflows: Multi-step approval processes

4.2.3 Context-Aware Processing

Project Association: Link documents to relevant projects

Customer Context: Associate with customer relationships

Historical Patterns: Learn from past routing decisions

Deadline Awareness: Consider time-sensitive requirements

Compliance Requirements: Route for regulatory compliance

4.3 Workflow Integration

4.3.1 Communication Integration

Slack Integration: Channel notifications and bot interactions

Microsoft Teams: Team collaboration and notifications

Email Notifications: Customizable email alerts

Mobile Push: Real-time mobile notifications

In-App Messaging: Internal messaging system

4.3.2 Task Management

Jira Integration: Automatic ticket creation and updates

Asana Integration: Task assignment and tracking
Trello Integration: Card creation and board management
Custom Webhooks: Integration with any REST API
Calendar Integration: Schedule document review sessions

4.4 Search & Discovery

4.4.1 Intelligent Search

Full-Text Search: Content-based document search
Semantic Search: Meaning-based search capabilities
Filter Options: Date, type, assignee, status filters
Faceted Search: Multiple filter combinations
Auto-Complete: Search suggestion and completion

4.4.2 Knowledge Discovery

Related Documents: Find similar or related content
Trend Analysis: Identify patterns in document types
Usage Analytics: Track document access and usage
Recommendation Engine: Suggest relevant documents
Knowledge Graphs: Visual relationship mapping

5. Technical Requirements

5.1 Performance Requirements

Processing Speed: <10 seconds per document classification
Throughput: 500+ documents per hour
Search Response: <1 second for search queries
Uptime: 99.5% service availability
Concurrent Users: Support 200+ simultaneous users

5.2 AI/ML Requirements

Classification Accuracy: >95% for standard document types
Model Training: Continuous learning from user feedback
Multi-Language: Support for 10+ major languages
Custom Models: Industry-specific model training
Confidence Thresholds: Configurable confidence levels

5.3 Integration Requirements

API Standards: RESTful APIs with comprehensive documentation
Real-time Events: WebSocket support for live updates
Webhook Support: Outbound event notifications
SSO Integration: SAML, OAuth 2.0 authentication
Data Export: Multiple format export capabilities

6. User Experience Design

6.1 Web Dashboard

Document Pipeline: Visual workflow status
Quick Actions: Common tasks and shortcuts
Analytics Dashboard: Usage and performance metrics
Configuration Panel: Rule and workflow management

6.2 Mobile Interface

Document Capture: Camera-based document upload
Quick Classification: Mobile document processing
Notification Center: All alerts and updates
Offline Capability: Basic functionality without internet

6.3 Admin Console

User Management: Role and permission administration
Rule Configuration: Business rule setup and management
Model Training: AI model management and tuning
Analytics & Reporting: Comprehensive usage analytics

7. Security & Compliance

7.1 Data Security

Encryption: End-to-end encryption for all documents
Access Control: Granular permission management
Audit Logging: Complete activity tracking
Data Residency: Configurable data storage locations
Backup & Recovery: Automated backup systems

7.2 Compliance Features

GDPR Compliance: Data protection and privacy controls

HIPAA Ready: Healthcare data handling capabilities

SOX Compliance: Financial document audit trails

Industry Standards: ISO 27001, SOC 2 compliance

Data Retention: Configurable retention policies

8. Success Metrics

8.1 Operational Metrics

Classification Accuracy: Percentage of correct classifications

Processing Time: Average time per document

User Adoption: Active users and feature usage

Search Success Rate: Successful search completion

8.2 Business Impact Metrics

Time Savings: Reduction in document handling time

Productivity Gains: Documents processed per user

Error Reduction: Decrease in misrouted documents

User Satisfaction: Net Promoter Score (NPS)

9. Implementation Roadmap

9.1 Phase 1 - Core Classification (2 Days POC)

Basic document upload and classification

Simple routing rules

OpenAI integration for content analysis

Basic notification system

9.2 Phase 2 - Enhanced Intelligence (Week 1-2)

Advanced classification models

Complex routing engine

Integration with Slack/Teams

Search capabilities

9.3 Phase 3 - Enterprise Features (Week 3-4)

Multi-tenant architecture

Advanced analytics

Custom model training

Compliance features

9.4 Phase 4 - AI Enhancement (Month 2)

Machine learning optimization

Predictive routing

Advanced insights

Knowledge graph features

10. Risk Assessment

10.1 Technical Risks

AI Accuracy: Variable classification performance

Scalability: High volume processing challenges

Integration Complexity: Multiple system integrations

Data Privacy: Sensitive document handling

10.2 Mitigation Strategies

Human-in-the-Loop: Manual review for low confidence

Microservice Architecture: Independent scaling

API-First Design: Standardized integrations

Privacy by Design: Built-in data protection

11. Competitive Advantages

11.1 Technical Differentiators

AI-First Approach: Native AI capabilities vs. bolt-on features

Real-time Processing: Immediate classification and routing

Context Awareness: Understanding document relationships

Learning System: Continuous improvement from usage

11.2 Business Differentiators

Rapid Deployment: 2-day POC to production in weeks

Cost Effective: 80% lower than enterprise solutions

User-Friendly: Intuitive interface requiring minimal training

Flexible Integration: Works with existing tool ecosystem

12. Conclusion

The Intelligent Document Classifier and Router addresses a critical productivity challenge in modern business

Technical Design Document: Intelligent Document Classifier & Router

1. System Architecture Overview

API Gateway

Classification Service

Routing Engine

Content Analysis

AI/ML Service

Workflow Integration

Search & Discovery

Shared Data Layer

Document Storage

Metadata DB

Permissions DB

2. Microservice Specifications

Service Technology Core Responsibilities Key Interfaces

API Gateway FastAPI, JWT Request routing, authentication, rate limiting REST endpoints, WebSocket

Classification Python, spaCy Document type detection, content categorization File upload API, confidence

Routing Engine Python, Rule Engine Context-aware assignment, workload balancing JSON rule definitions,

Content Analysis Pandas, NumPy Entity extraction, relationship mapping Dataframe processing, metadata

AI/ML Service scikit-learn Model training/inference, feedback integration Model versioning endpoint

Workflow Integration Python, Webhooks Jira/Slack connectivity, notification delivery Custom adapter patte

Search & Discovery SQLAlchemy, Full-Text Semantic search, knowledge graphs Query parser, ranking alg

3. Data Flow Sequence

4. Database Schema Design

Documents Table

Column Type Description

doc_id UUID Unique document identifier

original_name VARCHAR(255) Original filename

storage_path TEXT S3 path/object reference

doc_type VARCHAR(50) Contract/Invoice/Report etc.

confidence FLOAT Classification confidence score

Metadata Table

Column Type Description

meta_id SERIAL Auto-increment ID

doc_id UUID Foreign key to documents

key_entities JSONB Extracted names/dates/amounts

related_docs UUID[] Linked document references

risk_score FLOAT Compliance risk assessment

Routing Rules Table

Column Type Description

rule_id SERIAL Auto-increment ID

condition JSONB IF-THEN logic definition

assignee VARCHAR(100) User/team assignment target

priority INT 1-5 urgency scale

5. CI/CD Pipeline Design

6. Non-Functional Requirements

Category □ Requirement □ Implementation Strategy

Performance □ <10s/doc processing □ Async processing, message queues

Scalability □ 500 docs/hour throughput □ Kubernetes auto-scaling

Security □ End-to-end encryption □ TLS 1.3, AES-256 at rest

Reliability □ 99.5% uptime □ Health checks, circuit breakers

Maintainability □ 85% test coverage □ Pytest, mutation testing

Compliance □ GDPR/HIPAA readiness □ Audit logs, data anonymization hooks

8. Failure Mode Analysis

Failure Point □ Mitigation Strategy □ Recovery Process

OCR failure □ Fallback to metadata-based routing □ Human review queue

Classification low confidence □ Secondary model verification □ User feedback loop

Routing rule conflict □ Priority-based arbitration □ Dead-letter queue monitoring

Service downtime □ Kubernetes liveness probes □ Auto-container restart

Data corruption □ Cryptographic hash verification □ S3 version rollback

9. Monitoring & Observability

Key Metrics Dashboard

Metric □ Source □ Alert Threshold

Docs processed/min □ Prometheus □ <10

Classification accuracy □ Model feedback □ <90%

Avg. routing time □ Application logs □ >30s

Error rate □ ELK stack □ >5%

Container memory □ cAdvisor □ >80% utilization

10. Security Controls

Layer □ Controls

Network □ VPC isolation, Security Groups

Application □ JWT authentication, Input validation

Data □ AES-256 encryption, RBAC permissions

Infrastructure □ Immutable containers, Signed images

Compliance □ Audit trails, Data retention policies

This design enables implementation of core classification/routing functionality within 2 days using specified

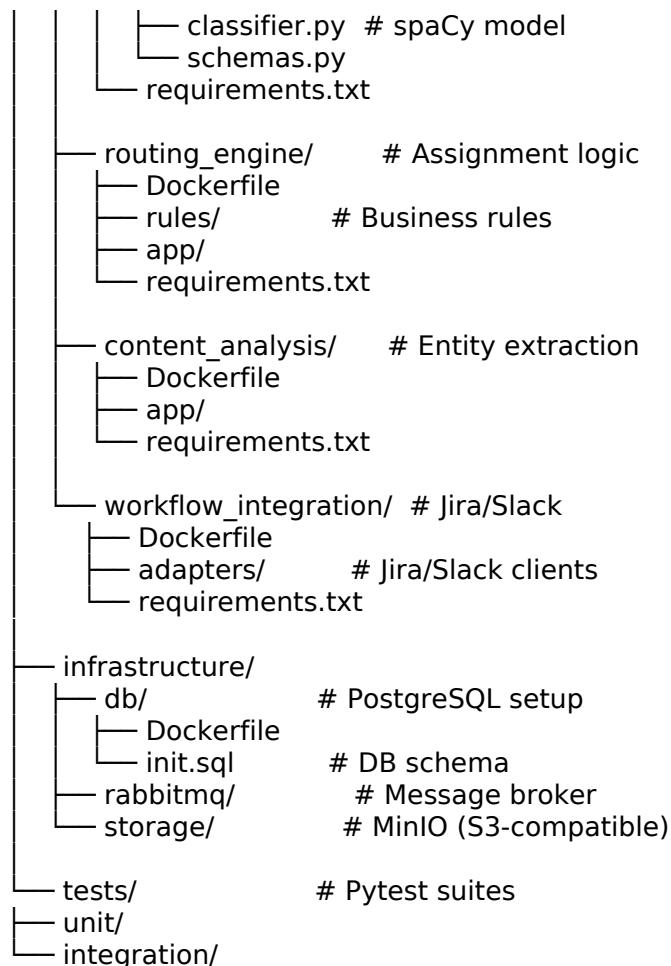
Initial Folder Structure & Docker Setup Guide

Tech Stack: Python, FastAPI, Docker Compose, PostgreSQL, SQLAlchemy, RabbitMQ, GitHub Actions

Root Directory Structure

intelligent_document_router/

```
├── .github/
│   └── workflows/
│       └── ci-cd.yml          # GitHub Actions pipeline
├── docker-compose.yml        # Main compose file
├── .env                      # Environment variables
├── docs/                    # Architecture diagrams
├── libs/                    # Shared Python libraries
│   ├── database/            # SQLAlchemy base models
│   └── utils/                # Common utilities
├── microservices/
│   ├── api_gateway/         # Entry point
│   │   ├── Dockerfile
│   │   ├── app/
│   │   │   ├── main.py      # FastAPI routes
│   │   │   └── dependencies.py
│   │   └── requirements.txt
│   └── classification/      # Doc classification
│       ├── Dockerfile
│       └── app/
```



Key Files Implementation

1. docker-compose.yml

version: '3.8'

services:

api_gateway:

build: ./microservices/api_gateway

ports:

- "8000:8000"

env_file: .env

depends_on:

- rabbitmq

- db

classification:

build: ./microservices/classification

env_file: .env

routing_engine:

build: ./microservices/routing_engine

env_file: .env

content_analysis:

build: ./microservices/content_analysis

env_file: .env

workflow_integration:

build: ./microservices/workflow_integration

env_file: .env

db:


```
build: ./infrastructure/db
env_file: .env
volumes:
- pgdata:/var/lib/postgresql/data
```

```
rabbitmq:
image: rabbitmq:3-management
ports:
- "5672:5672"
- "15672:15672"
```

```
minio:
image: minio/minio
command: server /data
volumes:
- storage:/data
ports:
- "9000:9000"
```

```
volumes:
pgdata:
storage:
2. Sample Dockerfile (for classification service)
FROM python:3.9-slim
```

```
WORKDIR /app
```

```
COPY ./requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
```

```
# Download spaCy model during build
RUN python -m spacy download en_core_web_sm
```

```
COPY ./app ./app
```

```
CMD ["uvicorn", "app.classifier:app", "--host", "0.0.0.0", "--port", "8001"]
```

```
3. .env File
```

```
# PostgreSQL
```

```
POSTGRES_HOST=db
POSTGRES_USER=admin
POSTGRES_PASSWORD=secret
POSTGRES_DB=document_db
```

```
# RabbitMQ
```

```
RABBITMQ_HOST=rabbitmq
RABBITMQ_QUEUE=document_queue
```

```
# MinIO Storage
```

```
MINIO_ENDPOINT=minio:9000
MINIO_ACCESS_KEY=minioadmin
MINIO_SECRET_KEY=minioadmin
```

```
4. Database Initialization (infrastructure/db/init.sql)
```

```
CREATE TABLE documents (
id UUID PRIMARY KEY,
filename VARCHAR(255) NOT NULL,
filepath VARCHAR(255) NOT NULL,
doc_type VARCHAR(50),
confidence FLOAT
);
```

```
CREATE TABLE routing_rules (
```

```
id SERIAL PRIMARY KEY,  
condition JSONB NOT NULL,  
assignee VARCHAR(100) NOT NULL  
);
```

Development Workflow

Initialize Services:

```
docker-compose up --build -d
```

Access Services:

API Gateway: <http://localhost:8000/docs>

RabbitMQ Console: <http://localhost:15672> (guest/guest)

MinIO Console: <http://localhost:9000> (minioadmin/minioadmin)

CI/CD Pipeline (.github/workflows/ci-cd.yml):

name: CI/CD Pipeline

on: [push]

jobs:

build-test:

runs-on: ubuntu-latest

steps:

- uses: actions/checkout@v3

- name: Build and test

run: |

```
docker-compose -f docker-compose.test.yml up --build --abort-on-container-exit
```

Microservice Communication Pattern:

```
# api_gateway/app/main.py
```

```
from fastapi import FastAPI, UploadFile
```

```
import httpx
```

```
app = FastAPI()
```

```
@app.post("/upload")
```

```
async def upload(file: UploadFile):
```

```
# 1. Save to MinIO
```

```
# 2. Send message to RabbitMQ
```

```
async with httpx.AsyncClient() as client:
```

```
await client.post("http://classification:8001/process", json=file_metadata)
```

Key Design Decisions

Decoupled Communication:

RabbitMQ for async inter-service messaging

HTTP for synchronous requests (service discovery via Docker DNS)

Shared Libraries:

Central libs/database with SQLAlchemy base models

Common validation schemas in libs/schemas.py

Stateless Services:

Each microservice connects directly to DB/MinIO

No inter-service dependencies during initialization

DevOps Integration:

Pre-configured Jenkinsfile in each service directory

GitHub Actions for automated testing

Docker Hub auto-builds on Git tags

POC-Scalable Architecture:

Local MinIO instead of AWS S3

SQLite option for lightweight testing

Mock AI models (spaCy) instead of cloud APIs

Next Steps for Implementation

Generate boilerplate code:

```
# For each microservice
mkdir -p microservice/{service}/app
touch microservice/{service}/Dockerfile
touch microservice/{service}/requirements.txt
```

Start core services:

```
docker-compose up db rabbitmq minio -d
```

Implement classification service first:

Document upload endpoint

spaCy text processing

DB storage logic

This structure allows parallel development by new hires, with each team owning one microservice while learning

Intelligent Document Classifier and Router - Implementation Todo

Project Overview

AI-powered document intelligence that automatically understands, categorizes, and routes business documents

Phase 1: Core Classification (2 Days POC) - PRIORITY

Day 1: Core Services Setup

1.1 Project Structure Creation

Create root directory structure

Set up microservices folders (api_gateway, classification, routing_engine, content_analysis, workflow_integrator)

Create shared libraries folder (libs)

Set up infrastructure folder (db, rabbitmq, storage)

Create tests folder structure

1.2 Docker & Infrastructure Setup

Create docker-compose.yml with all services

Set up PostgreSQL database with init.sql

Configure RabbitMQ message broker

Set up MinIO (S3-compatible storage)

Create .env file with all environment variables

Test infrastructure startup

1.3 Database Schema Implementation

Create documents table (UUID, filename, filepath, doc_type, confidence)

Create metadata table (meta_id, doc_id, key_entities, related_docs, risk_score)

Create routing_rules table (rule_id, condition, assignee, priority)

Set up SQLAlchemy models in libs/database

Test database connectivity

1.4 API Gateway Service

Create FastAPI main application

Implement JWT authentication

Set up request routing and rate limiting

Create document upload endpoint

Implement health check endpoints

Add API documentation (Swagger/OpenAPI)

1.5 Classification Service MVP

Set up spaCy for text processing

Implement document type detection (PDF, DOC, TXT, images)

Create basic classification logic (contracts, invoices, reports)

Add confidence scoring

Implement file storage to MinIO

Create classification API endpoints

Day 2: Integration & Deployment

2.1 Content Analysis Service

Implement text extraction from various formats

- Add OCR processing for images
- Create entity extraction (names, dates, amounts)
- Implement key phrase extraction
- Add summary generation
- Test with sample documents

2.2 Routing Engine

- Create rule engine for business logic
- Implement basic routing rules (role-based, expertise matching)
- Add workload balancing logic
- Create priority assignment system
- Implement routing decision tracking
- Test routing with sample scenarios

2.3 Workflow Integration

- Set up Slack integration (notifications)
- Implement Jira integration (ticket creation)
- Create email notification system
- Add webhook support for external systems
- Test notification delivery

2.4 Message Queue Integration

- Implement RabbitMQ producers/consumers
- Set up async document processing pipeline
- Create error handling and retry logic
- Add dead letter queue for failed messages
- Test message flow between services

2.5 Basic Search & Discovery

- Implement full-text search
- Add basic filtering (date, type, assignee)
- Create document relationship mapping
- Add simple analytics dashboard
- Test search functionality

2.6 POC Testing & Validation

- End-to-end testing with sample documents
- Performance testing (<10 seconds per document)
- Load testing (500+ documents per hour)
- User acceptance testing
- Document POC results and findings

Phase 2: Enhanced Intelligence (Week 1-2)

Week 1: Advanced Features

3.1 Advanced Classification Models

- Implement multi-language support
- Add industry-specific classifications
- Create custom category training
- Implement model fine-tuning
- Add confidence threshold configuration

3.2 Complex Routing Engine

- Implement conditional routing (if-then-else)
- Add multi-criteria routing logic
- Create escalation rules
- Implement approval workflows
- Add exception handling

3.3 Enhanced Content Analysis

- Add sentiment analysis
- Implement topic modeling

Create compliance scanning
Add risk assessment scoring
Implement relationship mapping
Week 2: Integration & Optimization
3.4 Advanced Integrations

Microsoft Teams integration
Asana/Trello integration
Calendar integration
Mobile push notifications
Custom webhook framework
3.5 Search & Discovery Enhancement

Implement semantic search
Add faceted search capabilities
Create knowledge graphs
Implement recommendation engine
Add trend analysis
Phase 3: Enterprise Features (Week 3-4)
Week 3: Multi-tenancy & Security
4.1 Multi-tenant Architecture

Implement tenant isolation
Add tenant-specific configurations
Create tenant management UI
Implement resource quotas
Add tenant analytics
4.2 Advanced Security

Implement end-to-end encryption
Add granular access control
Create audit logging system
Implement data residency controls
Add backup and recovery
Week 4: Compliance & Analytics
4.3 Compliance Features

GDPR compliance implementation
HIPAA readiness features
SOX compliance audit trails
Data retention policies
Compliance reporting
4.4 Advanced Analytics

Comprehensive usage analytics
Performance monitoring dashboard
Predictive analytics
Business intelligence reports
Custom reporting engine
Phase 4: AI Enhancement (Month 2)
Month 2: Machine Learning & Optimization
5.1 Machine Learning Optimization

Implement continuous learning
Add model versioning
Create A/B testing framework
Implement model performance monitoring
Add automated model retraining
5.2 Predictive Features

Predictive routing

Document priority prediction
Workload forecasting
Risk prediction
Trend forecasting
5.3 Advanced Insights

Document relationship discovery
Knowledge graph expansion
Intelligent recommendations
Process optimization suggestions
Business intelligence insights
Technical Implementation Details
Current Status: ☐ READY TO START
Product Requirements Document (plan1.md)
Technical Architecture Design
Implementation Roadmap
Project Structure Creation
Docker Setup
Database Schema
First Microservice
Next Immediate Actions:
Create the complete project folder structure
Set up docker-compose.yml with all services
Implement database schema and models
Create API Gateway service
Build Classification Service MVP
Success Criteria:
Document classification accuracy >95%
Processing time <10 seconds per document
Throughput 500+ documents per hour
99.5% service availability
Support 200+ concurrent users
Risk Mitigation:
Human-in-the-loop for low confidence classifications
Fallback mechanisms for service failures
Comprehensive error handling and logging
Performance monitoring and alerting
Security and compliance by design
Last Updated: [Current Date]
Current Phase: Phase 1 - Core Classification
Next Milestone: Day 1 - Core Services Setup

#-----

Intelligent Document Classifier and Router - Architecture & Folder Structure Guide

This guide explains the folder structure, code organization, and extension points for the Intelligent Document Classifier and Router.

☐ Root Directory Structure

```
intelligent_document_router/
├── .github/workflows/ # CI/CD pipeline definitions (GitHub Actions)
├── docs/              # Architecture diagrams, technical docs, PRDs
├── libs/              # Shared Python libraries (models, utils)
│   ├── database/      # SQLAlchemy models, DB connection logic
│   └── utils/          # Common utilities (logging, helpers)
├── microservices/     # All microservices (each in its own folder)
│   ├── api_gateway/   # FastAPI entry point, auth, routing
│   ├── classification/ # AI-powered document classification
│   ├── routing_engine/ # Smart document routing logic
│   ├── content_analysis/ # Entity extraction, content analysis
│   └── workflow_integration/ # Slack/Jira/email integrations
└── infrastructure/    # Infrastructure as code (DB, MQ, storage)
```

db/	# PostgreSQL Dockerfile, init.sql
rabbitmq/	# RabbitMQ config (if needed)
storage/	# MinIO/S3 config (if needed)
tests/	# Unit and integration tests
unit/	# Unit tests for each service
integration/	# Integration/E2E tests
docker-compose.yml	# Main compose file for all services
.env	# Environment variables for all services
README.md	# Project overview and quickstart
todo.md	# Implementation progress tracker
SETUP_COMPLETE.md	# Setup summary and next steps

□ Folder-by-Folder Explanation

.github/workflows/

Purpose: CI/CD pipeline definitions (e.g., build, test, deploy)

Add: YAML files for GitHub Actions, e.g., ci-cd.yml

docs/

Purpose: Architecture diagrams, PRDs, technical documentation

Add: Markdown docs, Mermaid diagrams, onboarding guides

libs/

Purpose: Shared code used by multiple microservices

database/` : SQLAlchemy models, DB session/connection logic

utils/` : Logging, error handling, helper functions

Add: New shared models, utility functions, or adapters

microservices/

Purpose: All business logic, each service in its own folder

api_gateway/` : FastAPI app, authentication, main API endpoints

app/` : Main FastAPI app, routers, middleware, dependencies

routers/` : Add new API endpoints (e.g., /documents, /auth)

middleware/` : Custom middleware (logging, rate limiting)

dependencies.py: Dependency injection (auth, DB session)

main.py: FastAPI app entry point

adapters/` : (Optional) HTTP clients for other services

rules/` : (Optional) Business rule definitions

requirements.txt, Dockerfile: Service dependencies and build

classification/` : Document type detection, spaCy/NLP logic

routing_engine/` : Rule engine, workload balancing, assignment

content_analysis/` : Entity extraction, key phrase, summary

workflow_integration/` : Slack, Jira, email, webhooks

Add: New microservices for additional features

infrastructure/

Purpose: Infrastructure as code for DB, MQ, storage

db/` : PostgreSQL Dockerfile, schema init.sql

rabbitmq/` : RabbitMQ config (if custom setup needed)

storage/` : MinIO/S3 config (if custom setup needed)

Add: Custom Dockerfiles, config scripts

tests/

Purpose: All automated tests

unit/` : Unit tests for each microservice

integration/` : Integration/E2E tests across services

Add: New test modules as you add features

□ Where to Add What Code

New API endpoint?

Add a new file in microservices/api_gateway/app/routers/ and register it in main.py.

New database model?

Add to libs/database/models.py and run migrations/init.

New microservice?

Create a new folder in microservices/, add app/, Dockerfile, requirements.txt.

Shared logic?

Add to libs/ (e.g., libs/utils/ for helpers, libs/database/ for DB logic).

Infrastructure change?

Update infrastructure/ and docker-compose.yml.

New test?

Add to tests/unit/ or tests/integration/ as appropriate.

□ Onboarding Guide for New Developers

Read README.md and ARCHITECTURE_GUIDE.md for project overview.

Clone the repo and copy .env.example to .env.

Start infrastructure:

docker-compose up db rabbitmq minio redis -d

Run the API Gateway locally:

cd microservices/api_gateway

pip install -r requirements.txt

uvicorn app.main:app --reload

Explore the codebase:

API endpoints: microservices/api_gateway/app/routers/

DB models: libs/database/models.py

Shared logic: libs/

Add new features/tests as described above

Run tests:

pytest

Check progress in todo.md and update as you complete tasks.

□ Extending the System

Add new document types: Update classification logic in classification/ and models in libs/database/models.py.

Add new routing rules: Update routing_engine/ and libs/database/models.py.

Integrate new tools (Slack, Jira, etc): Add adapters in workflow_integration/.

Improve search/analytics: Extend content_analysis/ and add endpoints to API Gateway.

□ Reference: plan1.md

All folder and service responsibilities are mapped directly from the PRD and technical design in plan1.md.

For detailed feature specs, see plan1.md and README.md.

For any questions, see the docs/ folder or contact the project maintainers.

□□ Example: Creating an API Endpoint in Each Microservice

Below are minimal examples for adding a simple /ping endpoint to each microservice. This helps new developers

1. API Gateway (microservices/api_gateway)

File: microservices/api_gateway/app/routers/ping.py

```
from fastapi import APIRouter
```

```
router = APIRouter()
```

```
@router.get("/ping")
```

```
def ping():
```

```
    return {"message": "pong from API Gateway"}
```

Register in app/main.py:

```
from app.routers import ping
```

```
app.include_router(ping.router, prefix="/ping", tags=["ping"])
```

2. Classification Service (microservices/classification)

File: microservices/classification/app/main.py

```
from fastapi import FastAPI
```

```
app = FastAPI()
```

```
@app.get("/ping")
```

```
def ping():
```

```
    return {"message": "pong from Classification Service"}
```

3. Routing Engine (microservices/routing_engine)

File: microservices/routing_engine/app/main.py

```
from fastapi import FastAPI
```

```
app = FastAPI()
```

```
@app.get("/ping")
```



```
def ping():  
    return {"message": "pong from Routing Engine"}  
4. Content Analysis (microservices/content_analysis)  
File: microservices/content_analysis/app/main.py
```

```
from fastapi import FastAPI  
app = FastAPI()
```

```
@app.get("/ping")  
def ping():  
    return {"message": "pong from Content Analysis Service"}  
5. Workflow Integration (microservices/workflow_integration)  
File: microservices/workflow_integration/app/main.py
```

```
from fastapi import FastAPI  
app = FastAPI()
```

```
@app.get("/ping")  
def ping():  
    return {"message": "pong from Workflow Integration Service"}  
How to test:
```

Start the relevant service (e.g., with `uvicorn app.main:app --reload` or via Docker Compose)
Visit `http://localhost:<service-port>/ping` in your browser or use `curl`
Tip:

For more complex APIs, create a `routers/` directory in each service and organize endpoints as in the API Gateway.
Always register new routers in your FastAPI `main.py` using `app.include_router(...)`.
Last changed by